

Systematic Hyperparameter Sensitivity Analysis of Random Forest: A Full Factorial Study on Classification and Regression Benchmarks

Pirapong Singsathid *

Abstract. A full factorial design of 1,024 Random Forest configurations, each evaluated by 10-fold cross-validation on 12 classification and 8 regression benchmarks, gives 204,800 model fits and an exact decomposition of performance variance, measured by eta-squared (η^2) with bootstrap intervals, Friedman tests and Shaffer-corrected Wilcoxon comparisons. Maximum tree depth dominates both task types, at mean $\eta^2 = 0.389$ for classification and 0.729 for regression, and is left unlimited in 17 of 20 best configurations. The number of features offered at each split matters proportionally more in classification (0.268) than regression (0.119), and is where tuning effort belongs: ensemble size and the minimum samples required to split a node can stay at their defaults, and depth should be left unconstrained rather than swept.

Keywords: Random Forest, hyperparameter sensitivity, full factorial design, eta-squared, Friedman test, Wilcoxon signed-rank test, benchmark evaluation, classification, regression

1 Introduction

Random Forest (RF) [1] consistently ranks among the most competitive off-the-shelf learning algorithms, yet its predictive performance depends on the choice of hyperparameters that govern tree construction and ensemble aggregation [10]. The cost of a poor choice is not marginal: across the benchmark used here, the spread between the worst and best of the 1,024 configurations reaches 0.375 in R^2 and 0.149 in weighted F1. This motivates systematic investigation of which parameters drive performance and by how much.

*Department of Mathematics, Faculty of Science, Khon Kaen University, Khon Kaen 40002, Thailand. E-mail: pirasi@kku.ac.th

Prior sensitivity studies share three limitations. They rely on local one-at-a-time designs or on surrogates fitted to observational records rather than on controlled experiment [13, 14]; they report which parameters matter, or whether tuning pays at all, rather than what share of performance variance each parameter explains [17]; and none covers regression.

This study addresses these gaps through four contributions:

1. A full factorial evaluation of all 1,024 RF hyperparameter configurations on 20 benchmark datasets spanning both classification and regression tasks, giving model-free estimates of every main effect and pairwise interaction.
2. A data-driven classification of all 20 datasets into three sensitivity regimes by dominant hyperparameter: `max_depth`-dominant, `max_features`-dominant, and `min_samples_leaf`-dominant, providing a basis for conditional tuning decisions.
3. Identification of systematic divergence between classification and regression sensitivity profiles: `max_depth` dominance is substantially larger in regression, a contrast invisible to classification-only benchmarks.
4. Evidence-based tuning recommendations separating the two parameters that re-pay tuning, `max_depth` and `max_features`, from the two that can be left at their defaults across every dataset tested, `n_estimators` and `min_samples_split`.

Section 2 reviews related work on RF hyperparameter analysis and sensitivity analysis methods. Section 3 describes the experimental methodology. Section 4 presents the results. Section 5 interprets the findings and their practical implications. Section 6 concludes the paper.

2 Related Work

Two things have to be in place before the experiment can be specified: a statement of how the hyperparameters enter Random Forest’s error, which fixes what a sensitivity measure should be measuring, and an account of what earlier empirical work has established about their relative importance, which fixes what this study has to add. The two subsections take them in that order.

2.1 Random Forest: Formal Model and Hyperparameter Structure

Random Forest [1] constructs an ensemble of T decision trees $\{h_1, \dots, h_T\}$, each grown on a bootstrap sample $\mathcal{B}_t$ drawn with replacement from training set $\mathcal{D}$ of n instances. At every internal node, the optimal split is selected over a randomly drawn subset of $m \leq p$ features, where p is the total number of input features. For a query point $\mathbf{x}$, the ensemble aggregates individual tree predictions through majority voting (classification) or averaging (regression):

$$\hat{y} = \arg \max_c \sum_{t=1}^T \mathbf{1}[h_t(\mathbf{x}) = c] \quad (1)$$

$$\hat{y} = \frac{1}{T} \sum_{t=1}^T h_t(\mathbf{x}) \quad (2)$$

For the regression case, the expected squared prediction error of the ensemble decomposes as [1]:

$$\mathbb{E}[(\hat{y} - y)^2] = \underbrace{\sigma_\varepsilon^2}_{\text{irreducible}} + \underbrace{\text{Bias}^2(\hat{y})}_{\substack{\text{max_depth} \\ \text{min_samples_leaf}}} + \underbrace{\rho_h \sigma_h^2}_{\substack{\text{max_features} \\ \text{max_depth}}} + \underbrace{\frac{1 - \rho_h}{T} \sigma_h^2}_{\text{n_estimators}} \quad (3)$$

where σ_ε^2 is the noise variance of the target, which no choice of hyperparameters can reduce, $\sigma_h^2 = \text{Var}(h_t(\mathbf{x}))$ is the variance of any single tree prediction, and $\rho_h = \text{Corr}(h_t(\mathbf{x}), h_{t'}(\mathbf{x}))$, $t \neq t'$, is the average pairwise inter-tree correlation. An analogous decomposition governs classification error [3].

None of the five hyperparameters appears in Eq. (3) directly. Each acts on one or more of its components, as the braces indicate, and the mechanism in each case is the following:

- **max_depth** (D) acts on two terms at once. A tree of depth D has at most 2^D leaves and can therefore return at most 2^D distinct values, which bounds how closely the ensemble can follow the response surface and sets Bias^2 ; deeper trees also fit their bootstrap sample more closely, which raises σ_h^2 . Because σ_h^2 scales both variance terms, depth reaches the total error through both.
- **min_samples_leaf** and **min_samples_split** block splits that would leave too few observations in a child node. This holds σ_h^2 down and raises Bias^2 , the same pair of channels as depth. Their effect is largely subsumed by **max_depth** when depth is constrained, and becomes the primary complexity control when it is not.
- **max_features** (m) acts on ρ_h . Drawing the candidate split features from a random subset of size $m < p$ forces different trees to split on different variables, and that decorrelation is what lowers ρ_h . Restricting m from p to $\lfloor \sqrt{p} \rfloor$ therefore reduces the correlated variance without directly inflating bias, so long as trees are deep enough to fit what the available features allow.
- **n_estimators** (T) enters only the uncorrelated variance term, as $1/T$. As $T \rightarrow \infty$ that term vanishes, but the correlated floor $\rho_h \sigma_h^2$ is untouched, so ensemble size yields strict diminishing returns.

The structure of Eq. (3) reveals an inherent asymmetry: beyond the irreducible σ_ε^2 , only ρ_h and σ_h^2 determine the error floor that persists as $T \rightarrow \infty$, while T itself provides a bounded benefit that saturates rapidly. This theoretical hierarchy motivates the empirical question at the center of this paper: how large are these effects in practice, and how do they vary across datasets and task types?

2.2 Prior Work and Research Gaps

Empirical investigations of RF hyperparameter importance have produced consistent but incomplete findings. Probst et al. [10] surveyed the influence of RF hyperparameters on predictive performance, identifying `max_features` and `min_samples_leaf` as commonly reported key parameters. A companion study by Probst et al. [11] quantified hyperparameter tunability across 38 algorithms and 38 OpenML datasets, finding that RF exhibits comparatively low global tunability relative to support vector machines, a statement about average behavior rather than about any particular parameter–dataset pair. van Rijn and Hutter [14] applied a functional ANOVA framework [6] to existing OpenML performance records, ranking `max_features` and `min_samples_leaf` as the dominant RF contributors. Mantovani et al. [7] covered 94 OpenML classification datasets, but for single-tree algorithms, which leaves the ensemble-specific parameters outside their scope. This line of work establishes which parameters matter on average; it does not measure how large each effect is.

Dedicated sensitivity analyses of RF hyperparameters have faced a methodological ceiling imposed by the tools they employed. Trithipkaiwanpon and Taetragool [13] applied one-at-a-time analysis and Latin Hypercube Sampling with ANOVA to four RF hyperparameters across four classification datasets, finding `min_samples_leaf` to be the most consistently sensitive parameter. OAT methods vary one parameter while holding the others at a chosen baseline, so they cannot detect interactions, and the effect they measure for any parameter depends on where the others were fixed [8]. The surrogate approach of van Rijn and Hutter [14] escapes that limitation but acquires another: a surrogate fitted to records in which `max_depth = None` is the prevailing default sees little variation in depth, and underestimates it accordingly. Section 5.1 returns to this point against the present estimates.

Large-scale empirical comparisons have framed the hyperparameter problem differently, asking whether tuning yields a worthwhile return rather than which parameters are responsible. Sipper [12] evaluated tuned versus default performance across 26 algorithms and 250 PMLB benchmark datasets using Bayesian optimization with 50 trials per run, finding that tuning provides modest average gains with substantial algorithm–dataset variation. Weerts et al. [17] introduced a tuning risk metric, the performance loss incurred when one hyperparameter is fixed at a default while the others are tuned, and showed across 59 OpenML-CC18 classification datasets that several can be fixed at computed defaults without measurable penalty. Both answer whether to tune; neither answers how much each parameter accounts for, and Weerts et al. leave interaction effects to future work.

The gap across these three lines of work is specific: no study combines a full factorial experimental design with standardized effect-size quantification across both classification and regression benchmarks. The cost of that omission is measurable in the present results: main effects account for 0.83 of performance variance in classification and 0.93 in regression, leaving 0.17 and 0.07 to interactions and residual variation, which one- at-a-time designs and surrogate rankings cannot attribute. The present study fills this gap by exhaustively evaluating all 1,024 hyperparameter configurations on 12 classification and 8 regression benchmark datasets, and by quantifying sensitiv-

ity through η^2 effect sizes, Friedman non-parametric tests, and Wilcoxon signed-rank post-hoc tests. The result is a complete, primary-evidence characterization of each hyperparameter’s role as predicted by Eq. (3).

3 Methodology

Equation (3) identifies $\text{Bias}^2(\hat{y})$, σ_h^2 , and ρ_h as the structural determinants of RF ensemble error, each modulated by a distinct subset of hyperparameters. The experiment below converts that characterisation into a measurable design.

3.1 Experimental Design

Let $\Theta = \{\theta_1, \dots, \theta_k\}$ denote the set of k hyperparameters under study, and let $L_i = \{\ell_{i,1}, \dots, \ell_{i,q}\}$ be the finite set of q discrete levels assigned to parameter θ_i . A full factorial design enumerates the complete Cartesian product

$$\mathcal{C} = L_1 \times L_2 \times \dots \times L_k, \quad |\mathcal{C}| = q^k \quad (4)$$

and evaluates every element of it independently. With $k = 5$ and $q = 4$ this yields $4^5 = 1,024$ configurations per dataset. Because every combination of levels is observed exactly once, all k main effects and all $\binom{k}{2} = 10$ pairwise interactions are estimable without assuming any functional form for the performance surface [8].

Selection of the parameter set. Four criteria determine which parameters enter Θ .

1. **Mechanistic coverage.** Equation (3) contains three hyperparameter-controllable quantities: the ensemble-size term $1/T$, the inter-tree correlation ρ_h , and the single-tree variance σ_h^2 together with Bias^2 . The set `n_estimators`, `max_features`, `max_depth`, `min_samples_leaf` and `min_samples_split` places each of these under experimental control, leaving no term unmanipulated and no two parameters acting solely through the same channel.
2. **Comparability.** These are the parameters ranked by the prior studies [10, 14, 17, 13] against which the present results are placed, so the η^2 shares reported below can be read directly against published importance rankings.
3. **Task-type invariance.** Each parameter carries an identical meaning and an identical admissible range in `RandomForestClassifier` and `RandomForestRegressor`, which is a precondition for the classification–regression comparison that motivates this study.
4. **Factorial tractability.** A complete design costs q^k evaluations per dataset, so every additional four-level parameter quadruples the experiment. At $k = 5$ the design requires 204,800 model fits; at $k = 6$ it would require 819,200, beyond

what exhaustive enumeration permits while retaining 20 datasets and 10-fold cross-validation.

Parameters deliberately excluded. Every other performance-relevant argument of the estimator is held at its scikit-learn default, each for a reason traceable to the criteria above. `bootstrap` is a structural switch rather than a tuning knob: disabling it removes the resampling mechanism that defines a Random Forest, so the resulting model is a different estimator rather than a differently tuned one. `max_samples`, `ccp_alpha` and `max_leaf_nodes` violate mechanistic coverage by redundancy. The first modulates ρ_h , the channel already controlled by `max_features`; the other two act on the σ_h^2 -Bias² channel already governed by `max_depth` and `min_samples_leaf`. Admitting them would partition one mechanism across more parameters rather than introduce a new one, deflating the measured η^2 of depth without altering the underlying effect. The split `criterion` violates task-type invariance: its levels are not comparable across the two settings (`gini`, `entropy` and `log_loss` for classification; `squared_error`, `absolute_error`, `friedman_mse` and `poisson` for regression), so it cannot enter a single factorial design shared by both. `class_weight` violates the same criterion, applying to classification alone, and modifies the training objective rather than the ensemble structure. What these exclusions imply for the interpretation of the reported effect sizes is stated in Section 5.4.

Table 1 lists the five hyperparameters, their four levels and the scikit-learn defaults. Two levels resolve only a linear contrast and three add a single curvature degree of freedom; neither can distinguish a saturating response from a monotone one, and saturation is precisely what Eq. (3) predicts for `n_estimators`. Four levels are the smallest grid on which saturation, non-monotonicity and an interior optimum are all observable, and the largest that keeps q^k within reach. Within each parameter the levels bracket the scikit-learn default and extend to the practical limit in the regularizing direction, so that the default is never an endpoint of the tested range. `max_features = 1.0` is included deliberately as the boundary case in which every feature is available at every split and ρ_h attains its maximum, fixing the upper endpoint of the correlation channel; `max_depth = None` plays the corresponding role for unconstrained tree growth.

Table 1: Hyperparameter grid: five parameters at four levels each, yielding 1,024 configurations.

Hyperparameter	Description	Levels	Default
<code>n_estimators</code>	Ensemble size T	50, 100, 200, 500	100
<code>max_depth</code>	Maximum tree depth D	None, 5, 10, 20	None
<code>max_features</code>	Features per split m	sqrt, log2, 0.5, 1.0	sqrt
<code>min_samples_split</code>	Min. samples to split	2, 5, 10, 20	2
<code>min_samples_leaf</code>	Min. samples at leaf	1, 2, 5, 10	1

All experiments use scikit-learn 1.7.2 [9] under Python 3.13, with `random_state = 42` and `n_jobs = -1` throughout; the statistical procedures of Section 3.3 use SciPy

1.16.3 [16].

3.2 Benchmark Datasets

Twenty benchmark datasets are used, 12 classification and 8 regression problems drawn from the scikit-learn built-in collection and from OpenML [15]. They were selected to vary the properties that plausibly moderate hyperparameter sensitivity rather than to maximize their number: sample size spans more than two orders of magnitude ($n = 150$ to $n = 21,613$), dimensionality ranges from $p = 4$ to $p = 64$, and the collection includes balanced and imbalanced targets, binary and many-class problems, and purely numeric as well as mixed numeric–categorical feature sets. All are publicly distributed, so the benchmark is reproducible without private data, and it overlaps the collections used by the studies against which the findings are compared in Section 5. The 12/8 split between task types was not imposed by design; η^2 , the Friedman test and the post-hoc comparisons are each computed within a dataset, so unequal group sizes affect the precision of the cross-dataset means reported in Section 4.1 but not the validity of the per-dataset estimates.

Tables 2 and 3 summarize the two groups. Categorical features in **credit-g** and **abalone** were encoded ordinally prior to model training, and a stratified random sample of 10,000 instances was drawn from **covertypes** to keep computation tractable.

Table 2: Classification benchmark datasets (12 datasets).

Dataset	OpenML ID	Samples n	Features p	Classes	Characteristics
Iris	—	150	4	3	Small, balanced
Wine	—	178	13	3	Small, balanced
Ionosphere	59	351	34	2	Binary, noisy features
Vehicle	54	846	18	4	Multi-class
Breast Cancer	—	569	30	2	Binary, medical
Credit-G	31	1,000	20	2	Imbalanced, categorical
Banknote	1462	1,372	4	2	Low-dimensional binary
Digits	—	1,797	64	10	High-dimensional, multi-class
Segment	36	2,310	19	7	Many classes
Phoneme	1489	5,404	5	2	Medium-scale binary
Covertypes	180	10,000 [†]	54	7	Large-scale, imbalanced
Pol	722	15,000	26	2	High-dimensional binary [‡]

— Source: sklearn built-in. [†]Stratified sample from 581,012 instances (random_state = 42). [‡]OpenML 722 is a binarized version of a regression dataset: the numeric target is split into two classes at its mean.

Table 3: Regression benchmark datasets (8 datasets).

Dataset	OpenML ID	Samples n	Features p	Characteristics
Diabetes	—	442	10	Small-scale, medical
Abalone	183	4,177	8	Categorical feature
Kin8nm	189	8,192	8	Nonlinear dynamics
Elevators	216	16,599	18	Large-scale
Bike Sharing	42712	17,389	12	Temporal pattern
Houses	537	20,640	8	Real estate pricing
California Housing	—	20,640	8	Geographic features
House Sales	42165	21,613	17	Mixed feature types

— Source: sklearn built-in.

3.3 Evaluation and Statistical Analysis

Cross-validation protocol. Let $\mathcal{D} = \{(\mathbf{x}_i, y_i)\}_{i=1}^n$ denote a dataset of n instances. Each configuration $\mathbf{c} \in \mathcal{C}$ is evaluated by 10-fold cross-validation: $\mathcal{D}$ is partitioned into folds $\{F_1, \dots, F_{10}\}$ of approximately equal size, and the model is trained on $\mathcal{D} \setminus F_j$ and evaluated on F_j for each j . Classification folds are stratified, so that each preserves the class proportions of the full dataset, i.e. $|F_j \cap \mathcal{D}_c|/|F_j| \approx n_c/n$ for every class c . A single partition was preferred to repeated cross-validation because the available model fits were allocated to breadth: r repetitions would multiply the cost by r , whereas the same number of fits buys complete enumeration of the configuration space on all 20 datasets. The cost of that choice is taken up in Section 5.4.

Performance metrics. For each configuration $\mathbf{c} \in \mathcal{C}$ and fold j , a scalar performance score $M_j(\mathbf{c})$ is recorded. For classification tasks, the primary metric is the weighted F1-score:

$$\text{F1}_w = \sum_{c=1}^C \frac{n_c}{n} \cdot \frac{2P_c R_c}{P_c + R_c} \quad (5)$$

where P_c and R_c are the per-class precision and recall, n_c is the number of true instances of class c , and $n = \sum_c n_c$. Weighted averaging accounts for class imbalance without discarding information from minority classes. For regression tasks, the primary metric is the coefficient of determination:

$$R^2 = 1 - \frac{\sum_{i=1}^n (y_i - \hat{y}_i)^2}{\sum_{i=1}^n (y_i - \bar{y})^2} \quad (6)$$

which is scale-free and therefore comparable across datasets with different target ranges. Accuracy and RMSE were also recorded for every configuration and are released with the per-dataset result files; they are not analysed here, because η^2

computed on them reproduces the ranking obtained from the primary metrics. For each configuration the mean $\bar{M}(\mathbf{c}) = \frac{1}{10} \sum_{j=1}^{10} M_j(\mathbf{c})$ and standard deviation $s_M(\mathbf{c})$ across folds are retained.

Eta-squared effect size. The sensitivity of hyperparameter θ_i on dataset $\mathcal{D}$ is quantified by the eta-squared (η^2) effect size, an ANOVA-style measure of the proportion of total performance variance attributable to variation in θ_i :

$$\eta^2(\theta_i) = \frac{SS_{\text{between}}(\theta_i)}{SS_{\text{total}}} = \frac{\sum_{\ell \in L_i} n_\ell (\bar{M}_\ell - \bar{M})^2}{\sum_{\mathbf{c} \in \mathcal{C}} (\bar{M}(\mathbf{c}) - \bar{M})^2} \quad (7)$$

where $\bar{M}_\ell$ is the mean performance over all configurations at level ℓ of θ_i , $n_\ell = |\mathcal{C}|/q$ is the number of such configurations, and $\bar{M}$ is the grand mean over all $|\mathcal{C}|$ configurations. The denominator SS_{total} includes variance contributions from all five hyperparameters, their interactions and residual variation, which makes η^2 a marginal measure: it answers what fraction of observable performance variance is attributable to θ_i alone. Being marginal, the five main effects need not sum to unity; over dataset means they account for 0.83 of total variance in classification and 0.93 in regression, the remainder being interaction and residual variance. η^2 is preferred here to the functional-ANOVA importance of van Rijn and Hutter [14] for a reason specific to the design: because every cell of the factorial is observed, the sum-of-squares decomposition in Eq. (7) is exact and needs no surrogate model to interpolate unobserved configurations, and it reports the result on a bounded $[0, 1]$ scale comparable across datasets and across the two performance metrics used here.

Uncertainty of the cross-dataset means. Each parameter’s headline η^2 is a mean over datasets, so the relevant sampling unit for its precision is the dataset. A 95% percentile bootstrap interval is obtained by resampling the per-dataset η^2 values with replacement within each task type over $B = 10,000$ resamples at seed 42, and taking the 2.5th and 97.5th percentiles of the resulting distribution of means. The interval quantifies sensitivity to the particular benchmark collection assembled here; it does not propagate uncertainty from the single cross-validation partition on which each η^2 is computed, which is a separate limitation discussed in Section 5.4.

Friedman test. For each hyperparameter θ_i and dataset $\mathcal{D}$, the Friedman test [4] assesses whether the $q = 4$ level-specific performance distributions differ significantly; this rank-based protocol, followed by pairwise post-hoc comparison, is the standard for comparing methods over multiple datasets [2]. Configurations are grouped by level and ranked within each group, and the test statistic

$$\chi_F^2 = \frac{12}{N_b q(q+1)} \sum_{j=1}^q R_j^2 - 3N_b(q+1) \quad (8)$$

follows a χ^2 distribution with $q - 1$ degrees of freedom under the null hypothesis of equal distributions, where R_j is the sum of ranks for level j and $N_b = |\mathcal{C}|/q$ is the number of configuration blocks. Significance is assessed at $\alpha = 0.05$.

Wilcoxon signed-rank post-hoc test. Where the Friedman test rejects the null hypothesis, all $\binom{q}{2} = 6$ pairwise level comparisons are conducted using the Wilcoxon signed-rank test [18], identifying the specific level contrasts that drive the observed effect. The six comparisons within a hyperparameter form one family, and family-wise error is controlled by Shaffer’s static procedure [5], which exploits the logical relations among pairwise hypotheses: with $q = 4$ levels the number of simultaneously true null hypotheses can only be 0, 1, 2, 3 or 6, these being the pair counts of the five partitions of four levels into equality classes. The ordered p -values are therefore compared step-down against $\alpha/6, \alpha/3, \alpha/3, \alpha/3, \alpha/2, \alpha$ rather than against the uniformly conservative $\alpha/6$ of a Bonferroni or Nemenyi correction. Table 9 reports both the uncorrected and the Shaffer-adjusted counts. The Bergmann–Hommel dynamic procedure is uniformly at least as powerful as Shaffer’s static procedure [5], so the adjusted counts are a lower bound on what that procedure would return.

4 Results

The subsections below move from the cross-dataset picture to the individual hyperparameters and then to their interactions, closing with the best configurations the grid contains. Effect sizes for every hyperparameter on every dataset are given in Tables 4 and 6, Friedman p -values in Tables 5 and 7, and the pairwise level contrasts in Tables 8 and 9. The text draws on those tables rather than restating them.

4.1 Global Sensitivity Profiles

Across all 20 benchmark datasets, `max_depth` is the dominant hyperparameter in both task types, `max_features` ranks second, and `min_samples_leaf` ranks third; `min_samples_split` and `n_estimators` are consistently negligible. This ranking is stable across task types but with markedly different magnitudes, as Figures 1 and 2 show. Bootstrap intervals on these means (Tables 4 and 6) are wide, as 12 and 8 datasets allow, and in classification the intervals for `max_depth` and `max_features` overlap. The separation that matters for the tuning recommendation survives nonetheless: neither leading interval reaches the classification upper bound for `min_samples_split`, 0.042, or for `n_estimators`, 0.008, and in regression `max_depth` is disjoint from every other parameter. Even the most favourable resampling of this benchmark leaves the `n_estimators` interval disjoint from those of both leading parameters, in both task types.

Pol warrants a robustness check, because it is not a native classification problem: OpenML 722 binarizes a numeric target at its mean, and its profile is correspondingly regression-like, carrying the third highest $\eta^2(\text{max_depth})$ of any classification dataset.

Excluding it lowers the classification mean for `max_depth` from 0.389 to 0.355 and shifts the other four means by less than 0.012, leaving the rank order unchanged. The task-type divergence of Section 4.3 does not depend on it either, and in fact widens without it, from 87% to 105%. Pol is retained in the reported figures, and every classification mean in this paper should be read as including one binarized regression problem.




Figure 1: Mean η^2 per hyperparameter, classification (blue) against regression (red), averaged over the 12 and 8 datasets of each task type.




Figure 2: Cross-dataset η^2 heatmap over all 20 benchmark datasets. Rows are hyperparameters ordered by mean η^2 ; columns are datasets, classification left and regression right. Darker red indicates higher sensitivity.

Classifying each dataset by its dominant hyperparameter yields three regimes. Thirteen datasets are **max_depth-dominant**: every regression dataset except Diabetes, together with Banknote, Digits, Segment, Phoneme, Covertype and Pol. Six are **max_features-dominant**: Iris, Wine, Ionosphere, Breast Cancer, Credit-G and Diabetes. Vehicle alone is **min_samples_leaf-dominant**. Figure 3 illustrates a representative example of each regime.




Figure 3: Hyperparameter sensitivity (η^2) for six representative datasets: California Housing and Diabetes (regression, contrasting depth sensitivity), Wine and Phoneme (classification, contrasting feature sensitivity), and Vehicle and Banknote (classification, contrasting leaf against depth dominance). The dotted line marks a conventional large-effect threshold, $\eta^2 = 0.14$.

The Friedman test rejects the null at $\alpha = 0.05$ for 97 of the 100 hyperparameter–dataset pairs. Each of the three remaining cases admits a dataset-specific explanation, given in Section 4.2: `n_estimators` on Ionosphere ($\chi_F^2 = 5.21$, $p = 0.163$), `max_features` on Phoneme ($\chi_F^2 = 0.43$, $p = 0.935$), and `max_depth` on Diabetes ($\chi_F^2 = 6.29$, $p = 0.108$).

4.2 Main Effects and Level Contrasts

max_depth. `max_depth` leads in both task types, and the gain it offers is concentrated almost entirely in one step. Depth-5 configurations are significantly worse than every deeper setting on 18 of 20 datasets (Shaffer-corrected; 19 uncorrected, Table 9), a near-universal performance floor produced by the 2^D leaf bound of Section 2.1. Contrasts among the deeper levels are far less consistent, 7 to 12 of 20, so the gain lies in escaping severe restriction rather than in choosing among deep settings; Figure 4 shows the corresponding steep rise from depth 5 to depth 10 and a plateau thereafter.

Sensitivity is highest on large classification datasets with complex decision boundaries, where Phoneme, Segment, Pol and Covertypes all exceed $\eta^2 = 0.73$, and it is near-universal in regression. Diabetes is the single exception ($p = 0.108$): with 442 samples and a low-complexity response surface, a depth constraint has little left to bind.

Table 4: Eta-squared (η^2) effect sizes for each hyperparameter across 12 classification datasets. Bold indicates the highest η^2 per row.

Dataset	n_est	max_d	max_f	mss	msl
Iris	0.0023	0.0000	0.2962	0.0081	0.2349
Wine	0.0030	0.0000	0.7985	0.0010	0.1061
Ionosphere	0.0021	0.0033	0.4915	0.0661	0.1930
Vehicle	0.0092	0.1842	0.0262	0.0405	0.4528
Breast Cancer	0.0211	0.0091	0.4099	0.1009	0.0821
Credit-G	0.0035	0.3287	0.4170	0.0036	0.1260
Banknote	0.0065	0.4356	0.0817	0.0257	0.3288
Digits	0.0047	0.4405	0.3649	0.0083	0.0683
Segment	0.0003	0.8727	0.0368	0.0030	0.0385
Phoneme	0.0015	0.8994	0.0040	0.0109	0.0356
Covertypes	0.0001	0.7328	0.1435	0.0071	0.0579
Pol	0.0000	0.7571	0.1473	0.0009	0.0137
Mean	0.0045	0.3886	0.2681	0.0230	0.1448
95% CI	0.002–0.008	0.200–0.588	0.146–0.403	0.008–0.042	0.078–0.222
Rank	5	1	2	4	3

n_est: n_estimators, max_d: max_depth, max_f: max_features,
mss: min_samples_split, msl: min_samples_leaf.

Table 5: Friedman test p -values for each hyperparameter across 12 classification datasets. Asterisk (*) denotes significance at $\alpha = 0.05$.

Dataset	n_est	max_d	max_f	mss	msl
Iris	<0.001*	0.007*	<0.001*	<0.001*	<0.001*
Wine	<0.001*	0.001*	<0.001*	<0.001*	<0.001*
Ionosphere	0.163	<0.001*	<0.001*	<0.001*	<0.001*
Vehicle	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Breast Cancer	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Credit-G	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Banknote	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Digits	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Segment	0.046*	<0.001*	<0.001*	<0.001*	<0.001*
Phoneme	<0.001*	<0.001*	0.935	<0.001*	<0.001*
Covertypes	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Pol	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*

n_est: n_estimators, max_d: max_depth, max_f: max_features,
mss: min_samples_split, msl: min_samples_leaf.

Table 6: Eta-squared (η^2) effect sizes for each hyperparameter across 8 regression datasets. Bold indicates the highest η^2 per row.

Dataset	n_est	max_d	max_f	mss	msl
Diabetes	0.0340	0.0100	0.3314	0.0355	0.2992
Abalone	0.0034	0.7881	0.0533	0.0015	0.0039
Kin8nm	0.0003	0.9340	0.0340	0.0063	0.0115
Elevators	0.0001	0.7925	0.1796	0.0021	0.0123
Bike Sharing	0.0002	0.8523	0.1274	0.0006	0.0046
Houses	0.0002	0.9582	0.0314	0.0004	0.0022
California Housing	0.0002	0.9778	0.0088	0.0004	0.0029
House Sales	0.0015	0.5205	0.1834	0.0175	0.1952
Mean	0.0050	0.7292	0.1187	0.0080	0.0665
95% CI	0.000–0.013	0.496–0.902	0.053–0.196	0.001–0.017	0.005–0.150
Rank	5	1	2	4	3

n_est: n_estimators, max_d: max_depth, max_f: max_features,
mss: min_samples_split, msl: min_samples_leaf.

Table 7: Friedman test p -values for each hyperparameter across 8 regression datasets. Asterisk (*) denotes significance at $\alpha = 0.05$.

Dataset	n_est	max_d	max_f	mss	msl
Diabetes	<0.001*	0.108	<0.001*	<0.001*	<0.001*
Abalone	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Kin8nm	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Elevators	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Bike Sharing	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
Houses	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
California Housing	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*
House Sales	<0.001*	<0.001*	<0.001*	<0.001*	<0.001*

n_est: n_estimators, max_d: max_depth, max_f: max_features,
mss: min_samples_split, msl: min_samples_leaf.




Figure 4: Main effect plots for selected classification (blue) and regression (red) datasets. Each panel gives the mean primary metric as a function of hyperparameter level, with shaded bands at ± 1 standard deviation across the 10 folds; yellow badges give η^2 for each parameter–dataset combination.

max_features. `max_features` ranks second in both task types, and the Wilcoxon contrasts give it a two-tier structure. `1.0` against `sqrt` is significant on 19 of 20 datasets, so full-feature trees are broadly inferior to the random subspace default; `log2` against `sqrt` is significant on only 9, so the two heuristics are largely interchangeable. Random subspace selection, not the particular fraction chosen, is the decisive mechanism. Table 8 shows the pattern on Wine, where `sqrt` and `log2` tie on every fold while all three restricted levels outperform `max_features = 1.0` at $p < 0.001$.

The effect is largest on smaller, low-dimensional classification datasets where depth is not the primary bottleneck: Wine, Ionosphere, Credit-G and Breast Cancer all exceed $\eta^2 = 0.40$. The non-significant case, Phoneme ($p = 0.935$), follows from its 5-feature space, where `sqrt` selects 2 features and `1.0` selects 5, a difference too small to affect ensemble diversity.

min_samples_leaf. `min_samples_leaf` ranks third, with its effect concentrated on small datasets where leaf-node overfitting is the primary failure mode: Vehicle, Banknote and Iris are the only classification datasets above $\eta^2 = 0.20$, and of the nine datasets with $n > 5,000$ only House Sales exceeds $\eta^2 = 0.06$. Vehicle is the sole `min_samples_leaf`-dominant dataset in the benchmark; with roughly 210 instances per class, raising the minimum leaf size from 1 to 2 or 5 supplies regularization that a depth constraint cannot replicate. The Wilcoxon counts place the effect at large leaf sizes: 20 of 20 datasets for 1 against 10, but 15 of 20 for 1 against 2.

min_samples_split and n_estimators. Both occupy ranks 4 and 5, with cross-dataset mean η^2 below 0.025, yet Wilcoxon detects significant level differences on most

datasets: 17 to 20 of 20 for the extreme `min_samples_split` contrasts and 13 to 17 for `n_estimators`. The differences are statistically real but practically negligible. Their main effect curves in Figure 4 are nearly flat throughout, and neither is a worthwhile tuning target under the conditions of this study.

4.3 Classification–Regression Divergence

The most consequential structural finding is the systematic divergence between the two task types. `max_depth` dominance is 87% larger in regression than in classification, while `max_features` runs the other way, mattering more in classification. Section 5.1 takes up the mechanism. The divergence is invisible to studies restricted to classification benchmarks.

4.4 Pairwise Interactions

The dominant pairwise interaction is between `max_depth` and `min_samples_leaf`. Figure 5 documents the depth effect across levels; Figure 6 maps the interaction structure for contrasting dataset–regime pairs.




Figure 5: Distribution of the primary metric across `max_depth` levels for four representative datasets; each box aggregates the 256 configurations sharing that depth value.

The distributions in Figure 5 separate the two regimes cleanly. At `max_depth` = 5 the interquartile range barely overlaps those at deeper settings, and the best configurations at depth 5 do not reach the median of unconstrained trees: severe depth restriction acts as a hard performance ceiling rather than a soft regularizer, so no combination of the remaining four parameters can compensate. At depth 10 and 20 the interquartile ranges overlap `max_depth` = `None` substantially, placing the saturation point between depth 10 and 20 for these datasets.

Figure 6 shows that the interaction structure is regime-specific rather than universal. California Housing varies almost entirely along the depth axis, from 0.632 at

Pairwise interaction heatmaps for four representative datasets




Figure 6: Pairwise interaction heatmaps for four representative datasets, each panel titled with the pair it shows. Top row: `max_depth` × `max_features` on California Housing (`max_depth`-dominant) and Wine (`max_features`-dominant). Bottom row: `max_depth` × `min_samples_leaf` on Vehicle (`min_samples_leaf`-dominant) and Diabetes (`max_depth` not significant). Cells give the mean primary metric over the 64 configurations sharing that level pair; the colour scale runs per panel from minimum (red) to maximum (green).

depth 5 to 0.809 unconstrained, while Wine varies almost entirely along the feature axis and is flat in depth: the dominant parameter sets the direction of the gradient, and the second shifts performance uniformly rather than conditionally. Interaction is therefore weak wherever a single parameter dominates.

Interaction becomes substantive once depth stops binding. On Vehicle the spread along `min_samples_leaf` is 0.012 in the `max_depth = 5` row and roughly 0.032 in each of the three deeper rows: a depth-5 tree is already too coarse for the leaf constraint to bite, and above that threshold leaf-level regularization becomes the primary remaining control on tree complexity. The change is a threshold at depth 5 rather than a gradual one, which is why `min_samples_leaf` attains its highest η^2 on datasets whose best configuration leaves `max_depth` unconstrained. Diabetes shows the same leaf gradient without the depth one: performance rises from 0.422 to 0.440 across `min_samples_leaf` while the depth axis is flat, consistent with its non-significant Friedman result for `max_depth`.

The same conditional logic applies to `max_features` on high-dimensional classification data such as Digits and Breast Cancer, where restricted feature subsets benefit most from unconstrained depth: deep trees compensate for the reduced feature set at each split, and pairing `max_features = 1.0` with shallow depth yields the weakest configurations.

Table 8: Wilcoxon signed-rank post-hoc test results for `max_features` on the Wine dataset (primary metric: weighted F1-score). Significant differences ($p < 0.05$) are indicated by *.

Level 1	Level 2	Statistic	p -value	Sig.
sqrt	log2	0.0	N/A [†]	—
sqrt	0.5	136.0	<0.001	*
sqrt	1.0	0.0	<0.001	*
log2	0.5	136.0	<0.001	*
log2	1.0	0.0	<0.001	*
0.5	1.0	0.0	<0.001	*

[†]All pairwise differences are zero (complete ties); p -value undefined.

4.5 Best Configurations Within the Tested Grid

Throughout this subsection, best denotes the highest-scoring configuration among the 1,024 evaluated, not a global optimum over the unrestricted parameter space. `max_depth = None` appears in 17 of 20 best configurations (Table 10), which establishes unconstrained tree growth as the default recommendation. The three exceptions, Ionosphere, Credit-G and House Sales, are all best at depth 10, and on none of them is `max_depth` the sole bottleneck: the first two are `max_features`-dominant, and House Sales carries the lowest `max_depth` η^2 of any regression dataset except Diabetes. Partial depth constraint thus acts as implicit regularization where the optimal

Table 9: Number of datasets (out of 20) in which each pairwise level comparison achieves Wilcoxon significance ($p < 0.05$). Raw counts are uncorrected; Shaffer counts apply Shaffer’s static procedure to the six comparisons within each hyperparameter at a family-wise $\alpha = 0.05$ (Section 3.3). Nine of the 30 contrasts lose one dataset under the correction; the other 21 are unchanged.

Hyperparameter	Level pair	Raw	Shaffer
max_depth	5 vs 10	19	18
	5 vs 20	19	18
	5 vs None	19	18
	10 vs 20	13	12
	10 vs None	12	12
	20 vs None	7	7
max_features	1.0 vs sqrt	19	19
	1.0 vs log2	18	18
	0.5 vs sqrt	17	17
	0.5 vs 1.0	16	15
	0.5 vs log2	16	16
	log2 vs sqrt	9	9
min_samples_leaf	1 vs 10	20	20
	2 vs 10	20	20
	2 vs 5	20	19
	5 vs 10	20	20
	1 vs 5	19	19
	1 vs 2	15	15
min_samples_split	2 vs 20	20	20
	10 vs 20	20	20
	5 vs 20	20	20
	2 vs 10	19	19
	5 vs 10	19	19
	2 vs 5	17	17
n_estimators	50 vs 100	17	16
	50 vs 200	17	17
	50 vs 500	17	17
	100 vs 500	15	14
	100 vs 200	14	13
	200 vs 500	13	13

decision boundary has limited complexity. `min_samples_leaf = 1` is optimal in 16 of 20 datasets, so minimal leaf regularization is preferable when depth is unconstrained. `n_estimators` ranges from 50 to 500 across the best configurations, yet its negligible η^2 means that variation carries little practical penalty: 100 to 200 trees suffice for most deployments.

For classification, best weighted F1 ranges from 0.7563 on Vehicle to 0.9942 on Banknote, median 0.9605; the three lowest scores, on Vehicle, Covertypes and Credit-G, all come from datasets with many classes, class imbalance or mixed-type features. For regression, best R^2 spans 0.4456 on Diabetes to 0.9466 on Bike Sharing, median 0.8267.

Table 10: Best hyperparameter configurations per dataset (ranked by primary metric on 10-fold CV), selected from the 1,024 configurations of the tested grid rather than from the unrestricted parameter space.

Dataset	Task	n_est	max_d	max_f	mss	msl	Metric	Score
Iris	Cls	100	None	1.0	5	1	F1	0.9597
Wine	Cls	50	None	sqrt	2	1	F1	0.9889
Ionosphere	Cls	50	10	sqrt	2	1	F1	0.9451
Vehicle	Cls	100	None	sqrt	5	2	F1	0.7563
Breast Cancer	Cls	50	None	0.5	5	1	F1	0.9612
Credit-G	Cls	500	10	0.5	10	1	F1	0.7659
Banknote	Cls	50	None	sqrt	2	2	F1	0.9942
Digits	Cls	200	None	log2	2	1	F1	0.9811
Segment	Cls	500	None	0.5	2	1	F1	0.9809
Phoneme	Cls	500	None	sqrt	2	1	F1	0.9137
Covertypes	Cls	500	None	0.5	2	1	F1	0.7610
Pol	Cls	100	None	0.5	2	1	F1	0.9869
Diabetes	Reg	50	None	0.5	2	10	R^2	0.4456
Abalone	Reg	500	None	0.5	20	5	R^2	0.5635
Kin8nm	Reg	500	None	1.0	2	1	R^2	0.7264
Elevators	Reg	500	None	1.0	2	1	R^2	0.8416
Bike Sharing	Reg	500	None	0.5	2	1	R^2	0.9466
Houses	Reg	500	None	0.5	2	1	R^2	0.8268
Calif. Hous.	Reg	500	None	log2	2	1	R^2	0.8266
House Sales	Reg	50	10	0.5	5	1	R^2	0.8701

`n_est`: n_estimators; `max_d`: max_depth; `max_f`: max_features; `mss`: min_samples_split; `msl`: min_samples_leaf. Cls: classification; Reg: regression. Score = weighted F1 (Cls) or R^2 (Reg).

5 Discussion

The results establish a consistent sensitivity hierarchy across 20 benchmark datasets, with `max_depth` and `max_features` accounting for the dominant share of performance variance in both task types.

5.1 Theoretical Interpretation of Sensitivity Rankings

The bias-variance-correlation decomposition of Eq. (3) accounts for the observed ranking. The five hyperparameters do not act on the same number of terms: as the braces in Section 2.1 show, `max_depth` alone governs two additive channels, while `max_features` and `n_estimators` each reach one. That asymmetry predicts a hierarchy with `max_depth` at the top, and the η^2 estimates match it: the dominance gap over the second-ranked parameter is 0.12 η^2 units in classification and 0.61 in regression.

For `n_estimators`, the $1/T$ factor in Eq. (3) means that increasing from 50 to 500 trees reduces the uncorrelated variance term by 90%, yet the correlated floor $\rho_h \sigma_h^2$ is entirely unaffected by ensemble size. Once $T \geq 100$, the uncorrelated term is already small relative to the correlated floor, making further increases negligible; the mean $\eta^2 < 0.005$ measures that ceiling. For `min_samples_split`, the redundancy with `min_samples_leaf` follows from the split feasibility constraint: a node can be split only if both resulting children retain at least `min_samples_leaf` observations, so a parent requires at least $2 \times \text{min_samples_leaf}$ samples irrespective of `min_samples_split`. Any setting with `min_samples_split` $\leq 2 \times \text{min_samples_leaf}$ is therefore inert, because the leaf constraint binds first. In the grid used here this covers 9 of the 16 `min_samples_split`–`min_samples_leaf` level combinations, leaving only 7 in which `min_samples_split` can alter tree growth at all. This explains its near-zero η^2 structurally, without reference to the datasets studied. Averaged over datasets the two parameters together account for 2.8% of total performance variance in classification and 1.3% in regression, though the classification figure conceals real spread: on Breast Cancer, where `min_samples_split` reaches $\eta^2 = 0.101$, the pair accounts for 12.2%.

The classification–regression divergence is consistent with target function geometry, on an interpretation this design does not test directly. A depth-5 tree offers at most 32 distinct constant predictions (Section 2.1), a severe constraint on a continuous response surface but not necessarily on a class label: a classification tree may still assign the majority of instances correctly across a complex boundary, because class boundaries are often lower-dimensional than the full response surface. Depth constraints should therefore cost more in regression, as the 87% gap in mean η^2 shows. The inverse pattern in `max_features` follows: once depth is unconstrained in regression, deep trees compensate for restricted feature sets at any given split, so `max_features` contributes less marginal gain.

This finding diverges from the conclusion of van Rijn and Hutter [14], who identified `max_features` and `min_samples_leaf` as the primary RF hyperparameters using functional ANOVA on OpenML meta-data. Two methodological differences account

for this divergence. First, their surrogate model was fitted on performance records in which `max_depth = None` is the scikit-learn default: when most training configurations already use unconstrained growth, the surrogate observes little variation in depth, systematically underestimating its sensitivity. Second, their benchmark comprised exclusively classification tasks from OpenML-CC18, omitting the regression settings where `max_depth` dominance is most pronounced. Primary experimental control removes both limitations here. The two sets of numbers remain measurements of different quantities, though: a surrogate importance fitted to observational records and an η^2 from a controlled enumeration answer one question under two designs, so the comparison is conceptual rather than head-to-head.

5.2 Dataset Moderators of Sensitivity

Two dataset characteristics systematically moderate hyperparameter sensitivity, both visible in the cross-dataset η^2 heatmap (Figure 2).

Sample size and max_depth sensitivity. `max_depth` η^2 increases with sample size (Figure 7). The mechanism is data-driven tree depth: with more observations, unconstrained trees grow deeper, widening the performance gap between `max_depth = None` and `max_depth = 5`. The trend is monotone within classification ($\rho = 0.942$), from Iris ($n = 150$, $\eta^2 \approx 0.000$) to Phoneme ($n = 5,404$, $\eta^2 = 0.899$), and class count is not the driver: Segment and Phoneme reach comparable sensitivity, $\eta^2 = 0.873$ and 0.899 , despite carrying 7 and 2 classes respectively.

In regression the size trend is present but statistically weak ($\rho = 0.419$, $p = 0.301$), for the opposite reason: every regression dataset with $n \geq 4,000$ already sits above $\eta^2 = 0.50$, leaving little room for size to discriminate. Diabetes (442 samples, $\eta^2 = 0.010$) is the only regression dataset below the ceiling, and California Housing (20,640 samples, $\eta^2 = 0.978$) sits at its top. Sample size is therefore an informative predictor in classification and a saturated one in regression.

Feature dimensionality and max_features sensitivity. The highest `max_features` η^2 values occur at intermediate dimensionality, between 10 and 34 features: Wine, Ionosphere and Breast Cancer all exceed $\eta^2 = 0.40$. Sensitivity falls away at both extremes. Phoneme and Banknote, with 5 and 4 features, drop below $\eta^2 = 0.09$ because `sqrt` and `1.0` converge to the same subset size. Digits and Covertypes, with 64 and 54 features, show only moderate sensitivity at $\eta^2 = 0.365$ and 0.144 , suggesting that information content per feature decreases as dimensionality grows, reducing the performance impact of any particular feature subset strategy.

5.3 Practical Tuning Recommendations

The η^2 effect sizes from this study convert the qualitative importance rankings of prior work [10, 14] into quantitative proportions, enabling practitioners to allocate tuning effort in proportion to measured variance.

1. **Leave unconstrained: max_depth.** `max_depth` carries the highest mean η^2 in




Figure 7: Scatter plot of $\max_depth \eta^2$ against dataset size (log scale) for all 20 benchmark datasets. Classification (circles, blue) and regression (squares, red) follow a broadly increasing trend (Spearman $\rho = 0.834$, $p < 0.001$ across all 20 datasets). The dashed line shows a linear trend fitted on the log-scale axis. Iris ($n = 150$, $\eta^2 \approx 0.000$) and Phoneme ($n = 5,404$, $\eta^2 = 0.899$) bracket the classification range; Diabetes ($n = 442$, $\eta^2 = 0.010$) and California Housing ($n = 20,640$, $\eta^2 = 0.978$) bracket the regression range.

both task types (Tables 4 and 6), but what follows from that is that the default is right, not that the parameter repays a sweep: `max_depth = None` is best on 17 of 20 datasets and `max_depth = 5` is significantly worse on 18 of 20 (Table 9). Leave depth unconstrained and spend no evaluations on it: on the six datasets with $n > 10,000$, `None` is both the scikit-learn default and the best-performing level, so a depth sweep there returns nothing. Constraining depth costs most in regression, where the mean η^2 reaches 0.729. Diabetes is the one dataset on which the choice is immaterial, the Friedman test not rejecting for `max_depth` ($p = 0.108$), and partial constraint pays only where the optimal boundary has limited complexity: Ionosphere, Credit-G and House Sales are best at depth 10.

2. **Tune first: `max_features`.** This is where the tuning gain sits. The `sqrt` default is best on only 5 of 20 datasets against 17 of 20 for the `max_depth` default, so a sweep here changes the selected configuration far more often. Priority is higher for classification and for datasets with more than 10 features. `log2` and `sqrt` are broadly interchangeable, significant on only 9 of 20 datasets, so a two-level sweep of `{sqrt, 1.0}` captures the dominant contrast at two evaluations.
3. **Tune second, conditionally: `min_samples_leaf`.** Relevant mainly on small datasets ($n < 1,000$), where it should be swept alongside `max_features` rather than after it, and wherever depth is left unconstrained, since leaf regularization is then the primary remaining complexity control. Its default of 1 is best on 16 of 20 datasets.
4. **Fix at defaults: `n_estimators` and `min_samples_split`.** Mean $\eta^2 < 0.025$ across all 20 datasets for both parameters: statistically detectable, practically negligible. Set `n_estimators` to 100–200 and `min_samples_split` to 2; tuning either parameter is unlikely to yield measurable performance gains on any dataset resembling those studied here.

5.4 Limitations

The following limitations qualify the results.

Restricted parameter set. The design manipulates five of the tuning arguments exposed by the implementation and holds the remainder at their defaults (Section 3.1). Every η^2 reported here is therefore a share of the variance generated within that five-parameter space, not within the full configuration space of the estimator. `ccp_alpha`, `max_leaf_nodes` and `max_samples` are expected to act through channels already represented, untested here, while `bootstrap` and `criterion` may carry effects a design admitting them would attribute away from the parameters studied. The ranking holds among the five parameters practitioners most commonly tune; it is not a claim that no other argument matters.

Single cross-validation replicate. Each configuration was evaluated once, under a 10-fold split generated from `random.state = 42`, so the design does not separate variance attributable to hyperparameters from variance attributable to that particular partition. Because η^2 is computed over 1,024 configurations sharing it, the partition

acts as a common factor rather than as independent per-configuration noise, and the large effects for `max_depth` are far too big to be an artefact of it. The near-zero values for `n_estimators` and `min_samples_split` need more caution, and the bootstrap intervals of Tables 4 and 6 do not supply it, since they resample across datasets rather than across partitions: a different partition would shift the very values they resample. Those values support the practical conclusion that tuning the two parameters is not worthwhile; establishing that their true effect is null would require repeated cross-validation over multiple seeds.

Discrete parameter grid and grid-induced contrast. The four levels per parameter do not cover the full continuous parameter space; effects at intermediate values (e.g., `max_depth` $\in \{7, 15\}$) are unobserved, and η^2 estimates reflect the grid used and may differ under finer resolution. In particular, the inclusion of `max_depth` = 5 creates a large performance gap to the next level (depth 10), which contributes to the high η^2 observed for `max_depth`. Had the grid begun at depth 10, the measured effect would be smaller. This is a known property of ANOVA-based sensitivity measures: effect size depends on the range of levels tested [8]. The grid spans the range a practitioner would plausibly consider, and a finer eight-level one would require $8^5 = 32,768$ configurations per dataset.

Ordinal encoding. Categorical features in datasets such as `credit-g` and `abalone` were encoded ordinally, which may introduce a spurious ordinal relationship. Tree-based methods are generally robust to monotone transformations, but target encoding or binary encoding could yield different sensitivity profiles for these datasets.

Covertypes subset. The stratified sample of 10,000 instances from 581,012 preserves class proportions but may not represent the full dataset’s spatial and feature distribution. Sensitivity results for Covertypes should be interpreted as applicable to the subset, not the full dataset.

scikit-learn implementation. The results are specific to `RandomForestClassifier` and `RandomForestRegressor` in scikit-learn 1.7.2. Alternative implementations (e.g., R `randomForest`, `ranger`) differ in default splitting rules and parallelization, which could affect the relative sensitivity of parameters such as `max_features`.

6 Conclusion

A full factorial design over 1,024 configurations, evaluated on 20 benchmark datasets, expresses RF hyperparameter importance as measured shares of performance variance. Four findings follow.

First, maximum tree depth dominates in both task types, and its default is what dominates: leaving depth unlimited is the best setting on 17 of 20 datasets, and capping it at five is significantly worse on 18 of 20. Second, each dataset falls into one of three sensitivity regimes named by its dominant parameter, depth-dominant (13 datasets), feature-subset-dominant (6) and leaf-size-dominant (1), and which regime applies is predicted in interpretable directions by sample size (Spearman $\rho = 0.834$

against the depth effect size, $p < 0.001$) and by feature dimensionality. Third, the profiles diverge by task type: depth dominance is 87% larger in regression than in classification, while the number of features per split matters proportionally more in classification ($\eta^2 = 0.268$ against 0.119). Fourth, the two leading parameters carry opposite practical consequences. The depth default should be left in place, whereas the default feature subset is best on only 5 of 20 and is therefore the setting that repays a sweep; ensemble size and the minimum samples required to split a node, at mean $\eta^2 < 0.025$ throughout, can be fixed at their defaults. All of this holds within the grid enumerated here, not over the unrestricted parameter space.

The same design applied to gradient boosting would answer the corresponding question there. XGBoost, LightGBM and CatBoost share tree depth, feature subsampling and leaf regularization with RF, but add a learning rate and boosting-specific parameters whose sensitivity profiles, and whose interactions with dataset characteristics, remain poorly characterized by primary experimental evidence.

Declarations

Acknowledgements. The author thanks Thapakorn Sinlapakorn for comments on the manuscript.

Funding. This research received no specific grant from any funding agency in the public, commercial, or not-for-profit sectors.

Conflict of interest. The author declares no conflict of interest.

Data and code availability. All benchmark datasets are publicly available through scikit-learn and OpenML [15]; the individual OpenML identifiers are listed in Tables 2 and 3. The experiment script and the complete per-dataset result files, statistical summaries and post-hoc test outputs are available under the MIT licence at <https://github.com/pirasi-kku/rf-hyperparameter-sensitivity>.

References

- [1] Breiman L., Random forests, *Machine Learning*, **45**, 1, 2001, 5–32.
- [2] Demšar J., Statistical comparisons of classifiers over multiple data sets, *Journal of Machine Learning Research*, **7**, 2006, 1–30.
- [3] Domingos P., A unified bias-variance decomposition and its applications, *Proceedings of the 17th International Conference on Machine Learning*, 2000, 231–238.
- [4] Friedman M., The use of ranks to avoid the assumption of normality implicit in the analysis of variance, *Journal of the American Statistical Association*, **32**, 200, 1937, 675–701.

- [5] García S., Herrera F., An extension on “Statistical comparisons of classifiers over multiple data sets” for all pairwise comparisons, *Journal of Machine Learning Research*, **9**, 2008, 2677–2694.
- [6] Hutter F., Hoos H., Leyton-Brown K., An efficient approach for assessing hyperparameter importance, *Proceedings of the 31st International Conference on Machine Learning (ICML)*, 2014, 754–762.
- [7] Mantovani R.G., Horváth T., Rossi A.L.D., Cerri R., Barbon Junior S., Vanschoren J., de Carvalho A.C.P.L.F., Better trees: An empirical study on hyperparameter tuning of classification decision tree induction algorithms, *Data Mining and Knowledge Discovery*, **38**, 3, 2024, 1364–1416.
- [8] Montgomery D.C., *Design and Analysis of Experiments*, 9th ed., John Wiley & Sons, 2017.
- [9] Pedregosa F., Varoquaux G., Gramfort A., Michel V., Thirion B., Grisel O., Blondel M., Prettenhofer P., Weiss R., Dubourg V., Vanderplas J., Passos A., Cournapeau D., Brucher M., Perrot M., Duchesnay É., Scikit-learn: Machine learning in Python, *Journal of Machine Learning Research*, **12**, 2011, 2825–2830.
- [10] Probst P., Wright M.N., Boulesteix A.-L., Hyperparameters and tuning strategies for random forest, *WIREs Data Mining and Knowledge Discovery*, **9**, 3, 2019, e1301.
- [11] Probst P., Boulesteix A.-L., Bischl B., Tunability: Importance of hyperparameters of machine learning algorithms, *Journal of Machine Learning Research*, **20**, 53, 2019, 1–32.
- [12] Sipper M., High per parameter: A large-scale study of hyperparameter tuning for machine learning algorithms, *Algorithms*, **15**, 9, 2022, 315.
- [13] Trithipkaiwanpon T., Taetrageel U., Sensitivity analysis of random forest hyperparameters, *Proceedings of the 18th International Conference on Electrical Engineering/Electronics, Computer, Telecommunications and Information Technology (ECTI-CON)*, 2021, 1163–1167.
- [14] van Rijn J.N., Hutter F., Hyperparameter importance across datasets, *Proceedings of the 24th ACM SIGKDD International Conference on Knowledge Discovery & Data Mining*, 2018, 2367–2376.
- [15] Vanschoren J., van Rijn J.N., Bischl B., Torgo L., OpenML: Networked science in machine learning, *ACM SIGKDD Explorations Newsletter*, **15**, 2, 2014, 49–60.
- [16] Virtanen P., Gommers R., Oliphant T.E., et al. and SciPy 1.0 Contributors, SciPy 1.0: Fundamental algorithms for scientific computing in Python, *Nature Methods*, **17**, 3, 2020, 261–272.
- [17] Weerts H.J.P., Müller A.C., Vanschoren J., Importance of tuning hyperparameters of machine learning algorithms, *arXiv preprint arXiv:2007.07588*, 2020.

- [18] Wilcoxon F., Individual comparisons by ranking methods, *Biometrics Bulletin*, **1**, 6, 1945, 80–83.